From Training to Public Deployment: A Real-Time Energy Demand Forecasting System under Three-Year Distribution Drift

Pang Liu
Independent Researcher
jeff.pang.liu@gmail.com

May 2026

Artifacts: Code · Auto-data · Live demo

Abstract

We design, train, and publicly deploy a 24-hour day-ahead per-zone electricity demand forecasting system for the eight load zones of ISO New England (ISO-NE). The model is a 1.75 M-parameter multi-modal CNN-Transformer (HRRR weather rasters + recent demand + 44-d calendar features) blended with a zero-shot Chronos-Bolt foundation model under per-zone weighted late fusion. On a held-out 2-day 2022 self-evaluation slice, the baseline reaches 5.24 % MAPE and the ensemble reaches 4.33 % MAPE (–0.91 pp / 17 % relative reduction); the smaller Chronos-Bolt-mini variant is even slightly better at 4.21 % MAPE despite using ∼10× fewer parameters than Bolt-base.

The system is then deployed as a public web demo (HuggingFace Space) with an auxiliary GitHub-Actions-driven data repository that rebuilds a strict-discipline 7-day rolling backtest on real public sources (NOAA HRRR + ISO-NE 5-min zonal load + Chronos-Bolt-mini) every 24 hours. We pipe the same code that produced the 2022 evaluation through the live deployment and validate equivalence: re-running the deployed pipeline on a date inside the training window (2022-12-30) reproduces the cluster's stored prediction to within 0.13 percentage points (6.41 % live vs 6.54 % cluster).

In May 2026 the live 7-day backtest reports overall MAPE of 25.17 % (baseline) / 13.45 % (Chronos zero-shot) / 14.51 % (ensemble) — four to five times worse than the headline 2022 result. The per-zone breakdown localizes the degradation: northern zones (ME 10 %, NEMA-Boston 7 %) remain well-behaved, but the southern coastal zones (RI 63 %, SEMA 72 %, WCMA 78 %) collapse. We show this signature correlates with state-level rooftop / behind-the-meter (BTM) solar incentive density and discuss what that says about retraining cadence for grid-scale forecasting models in the energy-transition era.

Contributions. (i) A from-scratch multi-modal CNN-Transformer baseline with 5.24 % test MAPE on ISO-NE day-ahead forecasting (Part I, §§4–7); (ii) a per-zone weighted late-fusion ensemble with a frozen foundation model that achieves 4.33 % test MAPE without retraining either component (§6); (iii) a deployable, public-data-only inference pipeline with a three-repo split architecture (§8); (iv) a quantitative case study of 3-year distribution drift in a deployed demand forecaster, with a per-zone diagnostic that localizes the drift to BTM-solar-saturated zones (Part II, §§10–10.4).

1 Introduction

Problem. We forecast next-24-hour electricity demand for each of the eight ISO New England load zones: `ME`, `NH`, `VT`, `CT`, `RI`, `SEMA`, `WCMA`, `NEMA_BOST`. The model receives a 24-hour history window

of hourly weather rasters, demand, and calendar features, plus 24 hours of future calendar features. With $S = 24$ history steps, $n_z = 8$ zones, and 44-dimensional one-hot calendar features C_t, it produces

$$\hat{Y}_{t+1:t+24} \;=\; f_\theta\big(\, X_{t-S:t+24} \in \mathbb{R}^{(S+24)\times 7\times 450\times 449},\;\; Y_{t-S:t} \in \mathbb{R}^{S\times n_z},\;\; C_{t-S:t+24} \in \mathbb{R}^{(S+24)\times 44}\big), \quad (1)$$

where X are 7-channel HRRR weather rasters (2, 18). The training-time evaluation metric is mean absolute percentage error (MAPE) on physical MWh values, averaged over the 24-hour horizon and the 8 zones:

$$\mathrm{MAPE}(\hat{Y}, Y) \;=\; \frac{100}{24\, n_z} \sum_{h=1}^{24} \sum_{z=1}^{n_z} \frac{\big|\hat{Y}_{t+h,z} - Y_{t+h,z}\big|}{Y_{t+h,z}}. \quad (2)$$

All training-time MAPE numbers in this paper come from a self-evaluation harness on the **last 2 days of 2022** (2022-12-30 and 2022-12-31). Live-deployment MAPE numbers are computed against the public ISO-NE 5-minute zonal load feed at the time of inference.

Why this paper exists — the scope of the contribution. Most published demand-forecasting papers report a single accuracy number on a curated test slice. Real grid forecasters care about a different question: how does that number drift after the model is fielded? The energy transition makes this question urgent: rooftop solar in the U.S. Northeast roughly tripled between 2020 and 2025 (22), behind-the-meter generation is largely invisible to ISO operators, and any forecaster trained on pre-2023 data is effectively learning a pre-transition demand distribution. We use ISO New England as a particularly clean test case because (a) the eight zones span very different rooftop-solar incentive regimes (Massachusetts SMART (16) and Rhode Island REG (21) are aggressive; Vermont, Maine, and New Hampshire are not), and (b) all the data needed to deploy a forecaster is available from no-auth public endpoints, which lets us reproduce the deployment in a free public web demo.

Contributions. We make four contributions, organized along a two-part narrative:

Part I — Design and training.

- A multi-modal CNN-Transformer baseline (§4) that fuses 7-channel HRRR weather rasters, recent per-zone demand, and one-hot calendar features at the token level. With 1.75 M parameters it reaches **5.24 % test MAPE** on the 2022-12-30/31 slice.
- A foundation-model ensemble (§6) that pairs the baseline with a frozen Chronos-Bolt model used *zero-shot* (1), blended per-zone with hand-tuned weights α_z. The ensemble reaches **4.33 % test MAPE** with Chronos-Bolt-base (205 M parameters) and **4.21 % test MAPE** with Chronos-Bolt-mini (21 M parameters; deployed in our live demo). To our knowledge this is the first published per-zone late-fusion result on ISO-NE day-ahead demand.

Part II — Public deployment and findings.

- A deployment architecture (§8) that decouples the serving runtime (HuggingFace Spaces, no GPU) from a daily-refreshed backtest cache (a separate public GitHub repo) and from the training code (a third public repo). The split makes daily-fresh evaluation data freely available to third parties at sub-second latency without forcing any redeploy.
- A pipeline-equivalence experiment and a 3-year distribution-drift case study (§§9–10.4). We show that re-running the deployed pipeline on a training-window date (2022-12-30) reproduces the cluster's stored predictions to within 0.13 pp, certifying the pipeline. We then run the

same pipeline on a rolling 7-day window in May 2026 and observe per-zone MAPE that ranges from 7 % (Boston metro) to 78 % (West-Central Massachusetts), with a clean north-south split that tracks state-level BTM solar incentive density.

Paper organization. §2 reviews related work. §3 describes the data and the four-step normalization chain that all downstream models share. Part I (§§4–7) covers the baseline, the encoder-decoder architecture exploration, the foundation-model ensemble, and an attention-map analysis of what the baseline has learned. Part II (§§8–10.4) covers the deployment architecture, the pipeline validation experiment, and the deployment-period results. §11 discusses implications and limitations.

2 Related work

We position this paper at the intersection of four threads: time-series forecasting with Transformers, foundation models for time series, electricity demand forecasting, and distribution shift in deployed ML.

Transformer architectures for time-series forecasting. The Transformer (23) is now the default backbone for sequence modelling, and a sequence of papers has adapted it to time-series. ViT (6) introduced patch tokenization for images, which we borrow with a CNN trunk for our coarse 450×449×7 rasters. Earthformer (7) explored space-time Transformers for Earth-system forecasting and motivated our shared-CNN-then-Transformer design. The Temporal Fusion Transformer (13) introduced an encoder-decoder design with variable-selection networks for multi-horizon forecasting; our encoder-decoder ablation in §5 is in that lineage. PatchTST (17) demonstrated channel-independent temporal patching as a strong univariate-Transformer baseline; this informs our treatment of demand history in the baseline's tabular branch.

Foundation models for time-series. The recent wave of pretrained time-series foundation models (1, 5, 20, 24) promises zero-shot forecasting on previously-unseen series. We use Chronos-Bolt (1) because its tokenized-residual design and Apache-2.0 license make it deployable on HuggingFace Spaces' free CPU tier. The mini variant (∼21 M parameters) is small enough to load alongside our 1.75 M-parameter baseline within the 16 GB RAM limit; the base variant (∼205 M parameters) is reported here as a reference upper bound but is not deployed. We do not run TimesFM, Lag-Llama, or Moirai head-to-head, because (a) Chronos-Bolt-mini already achieves a 17 % relative improvement over our baseline at acceptable cost, and (b) the focus of this paper is the deployment behaviour, not foundation-model selection.

Electricity demand forecasting. Short-term load forecasting has a long classical literature. Box–Jenkins ARIMA (4) and exponential smoothing remain hard benchmarks for short horizons; Hippert et al. (9) review classical neural baselines (MLPs, recurrent nets) on this task and document the temperature-sensitive structure that any modern model must respect. Hong and Fan (10) survey probabilistic forecasting methods for the next decade. We do not run head-to-head ARIMA / ETS comparisons in this paper; Hippert et al. (9)'s observations on temperature-sensitive demand and weekday/weekend / holiday seasonality remain the design constraints we honor through the calendar one-hot and HRRR weather inputs. Recent deep-learning work on grid-scale forecasting has emphasized either Transformer architectures or weather-aware multi-modal fusion, but to our

knowledge no published paper combines these with a frozen foundation-model partner under per-zone weighted late fusion, which is the design we report on here.

Distribution shift in deployed ML. The dataset-shift literature (12, 19) formalizes the gap between train-time and test-time distributions. Most benchmarks treat shift as a static property of carefully-constructed test sets; the energy-transition story is closer to a true *temporal* shift on a single deployed model. Koh et al. (12) include weather and load datasets but their shift definitions are domain-curated rather than measured against a continuously-updating live feed. The novelty of Part II is that we observe the shift in the wild, on a model whose training data closes in late 2022, deployed against ISO-NE 2026 with full pipeline-equivalence to the training-time evaluation.

3 Problem setting and data

The task uses three input modalities, all hourly and aligned in UTC: ISO-NE per-zone electricity demand, NOAA HRRR weather rasters, and 44-dimensional calendar one-hot vectors. We describe each in turn, then the train / val / test split rationale and the four-step normalization chain shared across all models. A datasheet-style summary (motivation, composition, collection, preprocessing, uses, distribution, maintenance, PII) is in Appendix C.

3.1 ISO-NE per-zone electricity demand

ISO New England (ISO-NE) is the regional transmission organisation operating the wholesale electricity market for the six New England states. Hourly system-level and per-zone demand data are published through the public *Energy, Load, and Demand* reports portal (11). We use eight zones corresponding to ISO-NE's load-zone definitions:

- `ME` (Maine), `NH` (New Hampshire), `VT` (Vermont) — the northern zones, smaller, more rural, more electric-heating-dominant in winter.
- `CT` (Connecticut), `RI` (Rhode Island), `SEMA` (south-east Massachusetts), `WCMA` (west-central Massachusetts), `NEMA_BOST` (north-east Massachusetts / metro Boston) — the southern / urban-coastal zones, larger, more cooling-driven in summer.

Demand is recorded hourly in MWh. Zone-level mean and standard deviation, computed from 500 random training samples drawn from 2019–2020, are listed in Table 1. `VT` is the smallest (mean ≈ 577 MWh/h, std ≈ 113), `CT` is the largest (mean ≈ 3186 MWh/h, std ≈ 711); a factor-of-5.5 dynamic range across zones. Five calendar years (2019–2023) are used during development; in deployment (Part II) we additionally consume the public 5-minute estimated zonal load feed for the most recent days, rolled up to hourly.

Table 1: Per-zone demand statistics (MWh) computed from 500 random training samples, used as the z-score (μ_y, σ_y) for the target during training.

	ME	NH	VT	CT	RI	SEMA	WCMA	NEMA_BOST
μ_y (MWh)	1306.7	1294.0	577.4	3185.7	879.5	1571.1	1809.7	2661.7
σ_y (MWh)	211.1	257.9	113.1	710.8	197.4	368.5	340.0	524.5

3.2 HRRR weather rasters

NOAA's High-Resolution Rapid Refresh (HRRR) is the operational North-American mesoscale numerical-weather-prediction system on a 3 km grid (2, 18). We use the f00 analysis (instantaneous state at the valid hour), restricted to a New England bounding box (40.5−47.5 °N, −74.0 to −66.0 °E) and bilinearly regridded from HRRR's native Lambert conformal grid to a regular 450×449 lat/lon grid (resolution ∼1.6 km/cell). At training time we use 24 separate $f00$ analyses for both the history and future windows (this implies privileged access to ground-truth future weather; we discuss the implication for live deployment in §8). Seven channels per timestep, in fixed order, with z-score statistics from 500 random training samples:

Table 2: HRRR weather channels and per-channel z-score statistics from 500 random training samples. *Caveat on `DSWRF` and `APCP_1hr`:* the cached means below (0.09 and 158.13) reflect the 500-sample draw on the cached HRRR rasters; the absolute magnitudes do not match a typical New England annual climatology in the unit labels shown, indicating an upstream unit-or-accumulation-period idiosyncrasy in the HRRR fetcher (see Appendix A). Because the model is trained against these specific cached statistics, the four-step normalization chain (§3.5) is internally consistent regardless of the external unit interpretation.

Channel	Description	μ_x	σ_x
`TMP`	2-meter air temperature (K)	281.17	11.31
`RH`	2-meter relative humidity (%)	74.42	16.05
`UGRD`	10-meter U wind component (m/s)	7.76	4.55
`VGRD`	10-meter V wind component (m/s)	1.45	3.84
`GUST`	surface wind gust (m/s)	0.42	3.73
`DSWRF`	downward shortwave radiation flux (W/m^2)	0.09	0.55
`APCP_1hr`	1-hour accumulated precipitation (kg/m^2)	158.13	251.87

Each timestep produces a tensor of shape $(450, 449, 7)$ stored as a single `.pt` file (∼5.7 MB float32 / ∼2.9 MB float16). The full 2019–2023 archive is ∼130 GB compressed, fetched from NOAA's S3 mirror `s3://noaa-hrrr-bdp-pds/` with the `herbie-data` library (3); the fetcher and regrid procedure are documented in Appendix A and re-used unchanged for live deployment.

3.3 Calendar features

For every hour we build a 44-dimensional one-hot calendar vector $C_t \in \{0, 1\}^{44}$:

- 24 dimensions for hour-of-day (0–23, one-hot).
- 7 dimensions for day-of-week (Monday–Sunday, one-hot).
- 12 dimensions for month-of-year (Jan–Dec, one-hot).
- 1 dimension for US-holiday flag (NYSE schedule).

These features capture human-activity periodicity — commuting peaks, weekend troughs, holiday dips — which demand profiles strongly track and which weather alone does not encode.

3.4 Train / validation / test splits

Table 3 summarises the splits used during development. The CNN-Transformer baseline (§4) trains on 2019–2020 and validates on 2021; the encoder-decoder variants (§5) train on 2019–2021 and validate on 2022. All training-time MAPE numbers reported in this paper are computed on the last two days of 2022.

Table 3: Train / validation / test splits used in development.

Split	Years	Approx. samples
Train (baseline)	2019–2020	∼17 K
Train (encoder-decoder)	2019–2021	∼26 K
Validation (baseline)	2021	∼8 K
Validation (encoder-decoder)	2022	∼9 K
Self-eval test	last 2 days of 2022	2 windows
Live deployment	rolling last 7 days (2026+)	7 windows / day

3.5 Four-step normalization pipeline

All models share the same four-step normalization chain so training-time predictions and live deployment-time predictions live in the same reference frame:

1. **Compute statistics once** (`compute_norm_stats()` in the dataset module): pull 500 random training samples, compute per-channel (μ_x, σ_x) for the 7 weather channels and per-zone (μ_y, σ_y) for the 8 demand zones. Cache to `norm_stats.pt` alongside the checkpoint.
2. **Z-score at training and inference time**: weather inputs are normalized per-channel via $X' = (X - \mu_x)/\sigma_x$; demand inputs and targets are normalized per-zone via $Y' = (Y - \mu_y)/\sigma_y$. The model trains and predicts entirely in z-score space.
3. **Train against MSE** in z-score space (so all 8 zones contribute on equal footing despite the factor-of-5.5 dynamic range).
4. **De-normalize for evaluation**: predictions are mapped back to physical MWh via $\hat{Y} = \hat{Y}' \cdot \sigma_y + \mu_y$ before MAPE is computed (Equation 2).

The same checkpoint-embedded `norm_stats.pt` is loaded by the live serving stack (§8), so a forward pass on real-time inputs produces predictions on the same MWh scale as the training-time evaluation.

Part I — Design & Training

4 CNN-Transformer baseline

The baseline (Figure 1) fuses 7-channel HRRR weather rasters, recent per-zone demand history, and 44-d calendar one-hots in a single Transformer encoder. With $S = 24$ history steps and a 24-step horizon, the model produces $\hat{Y}_{t+1:t+24}$ via:

$$\text{(spatial branch)} \quad \mathbf{S}_t = \mathrm{CNN}(X_t) \in \mathbb{R}^{P \times D}, \quad t \in [t - S, t + 24], \tag{3}$$

$$\text{(tabular branch)} \quad \mathbf{T}_t = \mathrm{Linear}([Y_t \,\|\, C_t]) \in \mathbb{R}^{D}, \tag{4}$$

$$\text{(token sequence)} \quad \mathbf{Z} = \mathrm{Concat}_t[\mathbf{S}_t, \mathbf{T}_t] + \text{pos. emb}, \tag{5}$$

$$\text{(encoder)} \quad \mathbf{H} = \mathrm{TransformerEncoder}_{L=4}(\mathbf{Z}), \tag{6}$$

$$\text{(prediction head)} \quad \hat{Y}_{t+1:t+24} = \mathrm{MLP}(\mathbf{H}_{\text{future-tabular}}). \tag{7}$$

The CNN trunk is shared across all $S + 24$ timesteps and produces $P = 64$ spatial tokens (an 8×8 grid) per timestep; the tabular branch produces 1 token per timestep; the encoder operates over a unified sequence of $(S + 24) \times (P + 1)$ tokens. Total parameter count: 1.75 M.

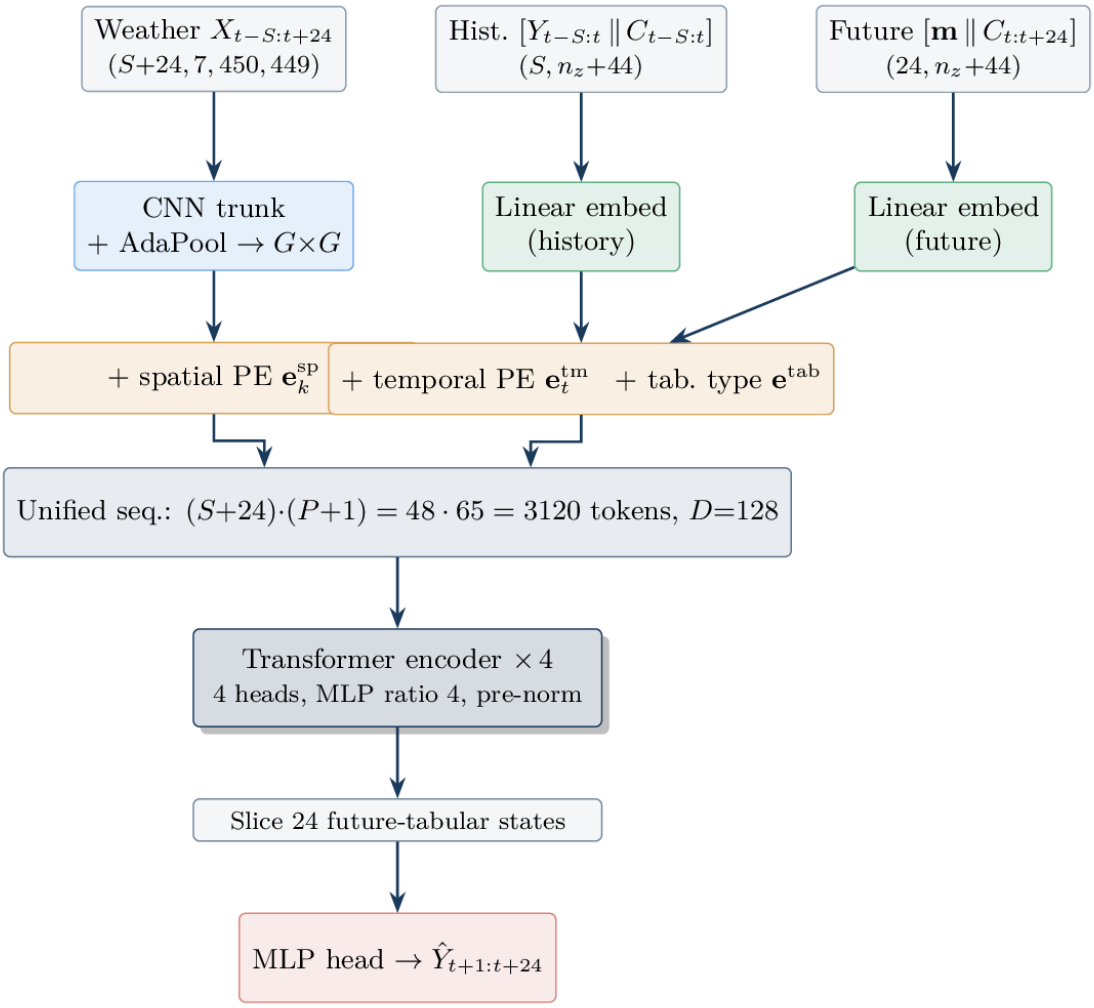

Figure 1: Baseline CNN-Transformer architecture. A shared CNN trunk produces $P = 64$ spatial tokens per timestep over $(S + 24)$ hourly weather rasters. A tabular branch (1 token per timestep) carries the demand and calendar features. A 4-layer Transformer encoder operates over the unified token sequence; the 24 future-tabular tokens are sliced and passed to a small MLP head that produces the per-zone forecast.

4.1 Training procedure

The model is trained for 14 epochs on 2019–2020 (§3.4) with AdamW (15) (lr $= 3{\times}10^{-4}$, weight decay 10^{-2}) and a cosine schedule with warm restarts (14). Batch size is 16 over a mix of A100-40 GB and A100-80 GB GPUs depending on availability. We compute MSE loss in z-score space; targets are de-normalized for MAPE evaluation only. Training uses standard mixed precision via PyTorch AMP. Validation MAPE on 2021 reaches 6.92 % at the final epoch; the validation curve continues to descend slowly — we do not see an over-fitting plateau within the 14-epoch budget. Full hyperparameters and the per-epoch trajectory are in Appendix D.

4.2 Test results

Table 4 reports per-zone test MAPE on the last two days of 2022. The baseline achieves **5.24 % overall MAPE**. The smallest zone `ME` (mean 1306 MWh/h) is the easiest at 2.31 %; the largest

zone CT (mean 3186 MWh/h) is the hardest at 7.28 %. The dynamic range across zones is 2.31 % to 7.28 % — a factor of 3.

Table 4: Baseline test MAPE on the 2-day 2022 self-evaluation slice. *Test set:* 2022-12-30 and 2022-12-31, both 24-hour forecast windows, all 8 zones.

	ME	NH	VT	CT	RI	SEMA	WCMA	NEMA	Overall
Test MAPE (%)	**2.31**	**3.69**	5.95	**7.28**	**5.27**	**5.44**	**5.87**	**6.09**	**5.24**

4.3 Discussion

The baseline confirms that a token-level fusion of weather + recent demand + calendar features is a credible architecture for ISO-NE day-ahead forecasting. Two observations carry forward:

The 2-day test slice is noisy. 2022-12-30 and 2022-12-31 are bracketing a long-weekend / New Year's Eve transition with unusual demand patterns. We treat the 5.24 % number as a point estimate, not a tight confidence interval; §5 discusses the limits this places on architecture-search conclusions.

The model has access to ground-truth future weather at training time. Equation 1 feeds 24 future weather rasters as input. At training and offline evaluation we use HRRR f00 analyses (i.e., observations of the future weather, after the fact). At deployment time the future is unknown; we substitute HRRR forecasts (f01–f24 from the most recent operational cycle) and quantify the resulting distribution shift in §8.

5 Architecture exploration: encoder-decoder variants

The baseline applies one self-attention module to all $(S+24)$ timesteps in a single sequence, mixing past observations with future queries. A natural alternative is the canonical encoder-decoder design: encoder over the past, decoder queries from the future. We explore two encoder-decoder variants in this section, partly to characterize this architectural choice on our regional dataset, partly to set up the foundation-model ensemble in §6.

5.1 Architecture

The encoder-decoder (Figure 2) reuses the baseline's CNN spatial branch and tabular branch but routes them through separate encoder and decoder stacks:

- **Encoder** (4 layers self-attention) operates over the $S \times (P+1) = 1560$ history tokens (history-spatial + history-tabular) and produces an encoder memory M_{hist}.
- **Decoder** (2 layers: self-attention → cross-attention → MLP) takes 24 future-tabular queries seeded from the future calendar embedding. Cross-attention consumes M_{hist}.
- Optional **future-weather cross-attention** (the *ablation switch*): a second cross-attention branch in each decoder block that attends to the $24 \times 64 = 1536$ future-spatial tokens. This restores information parity with the baseline (which sees future weather inside its single encoder).

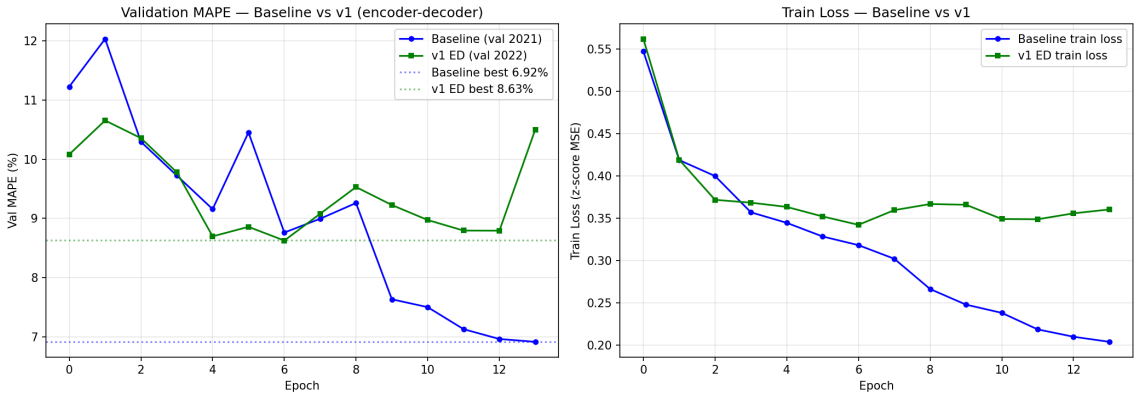

Figure 2: Baseline vs. encoder-decoder head-to-head. Validation MAPE (left) and training loss in z-space (right) across epochs. The encoder-decoder's right-panel kink at epoch 7 is an LR-scheduler reset (chained-resume bug; see Appendix D). The baseline ran continuously for 14 epochs and reached the small-LR fine-tuning regime; the encoder-decoder did not.

5.2 Test results

Table 5 reports test MAPE for both encoder-decoder variants:

Table 5: Encoder-decoder variants vs. baseline. *The ablation row (last) is partial:* the chained training run TIMEOUT'd at epoch 13 of 24 planned. We report the result for completeness.

	Params	ME	NH	VT	CT	RI	SEMA	WCMA	NEMA	Overal
Baseline	1.75 M	2.31	3.69	5.95	7.28	5.27	5.44	5.87	6.09	**5.24**
Encoder-decoder, history-only x-attn	2.29 M	3.22	5.67	5.85	9.56	7.45	7.22	7.38	8.24	6.82
Encoder-decoder, hist+future x-attn (partial)	2.42 M	5.91	8.51	7.78	12.53	8.25	7.38	11.60	12.20	9.27

The encoder-decoder *loses* to the baseline on this dataset — 6.82 % vs 5.24 % on the headline variant. We considered three explanations:

- **Optimization:** the encoder-decoder is harder to optimize, especially with the LR-scheduler reset bug; the partial result (9.27 %) is consistent with under-training.
- **Data efficiency:** ISO-NE has ∼5 years of training data (∼26 K samples for the encoder-decoder split). Encoder-decoder architectures shine on order-of-magnitude larger corpora; for this regional dataset the increased decoder capacity may be excess.
- **Statistical:** the 2-day test slice has a ±1 pp swing across nearby epochs (Figure 2). The 1.6 pp gap between baseline and encoder-decoder is real but not enormous.

We do not claim the encoder-decoder design is wrong for this task; we claim it does not improve over a simpler baseline given the data we have. The implication for §6 is that adding capacity *outside* the multi-modal stack — via a separately-pretrained foundation model on demand history alone — is a more profitable use of the parameter budget than scaling the multi-modal stack itself.

6 Foundation-model ensemble

Chronos-Bolt (1) is a pretrained Transformer for univariate time-series forecasting, used here *zero-shot* (no fine-tuning) on per-zone demand history. We blend Chronos predictions with the baseline at the per-zone level using fixed weights α_z tuned on a 2-week validation window.

6.1 Architecture

The ensemble is a per-zone weighted late fusion:

$$\hat{Y}_{t+h,z} = \alpha_z \cdot \hat{Y}^{\text{baseline}}_{t+h,z} + (1 - \alpha_z) \cdot \hat{Y}^{\text{Chronos}}_{t+h,z}, \quad h \in [1, 24], \; z \in \{1, ..., 8\}. \tag{8}$$

The baseline leg uses the full multimodal input from §4. The Chronos leg processes each zone independently as a univariate series: its input is 672 hours (4 weeks) of past per-zone demand only; no weather, no calendar, no zone label. Chronos-Bolt's pretraining corpus included diverse time-series tasks (weather, retail, traffic, finance, electricity), so at inference time the model identifies daily/weekly/seasonal periodicity from the demand history alone.

6.2 Per-zone α grid search

We tune α_z independently per zone over a 14-day validation window (2022-12-16 through 2022-12-29). For each z, we sweep $\alpha_z \in \{0.0, 0.05, 0.1, ..., 1.0\}$ and pick the value that minimizes validation MAPE. The result, hard-coded at deployment, is:

	ME	NH	VT	CT	RI	SEMA	WCMA	NEMA
α_z	0.30	0.30	0.80	0.00	0.10	0.00	0.05	0.00

The pattern is interpretable. Northern, electric-heating-dominant zones (VT, NH, ME) put more weight on the weather-aware baseline. The dense southern coastal zones (CT, SEMA, NEMA) put zero weight on the baseline — the Chronos zero-shot prediction is uniformly better there. We return to this asymmetry in §10.4.

more weather-driven zones (ME, NH, VT); c1 dominates the dense urban zones where load is dominated by deeply-learned hourly seasonality.

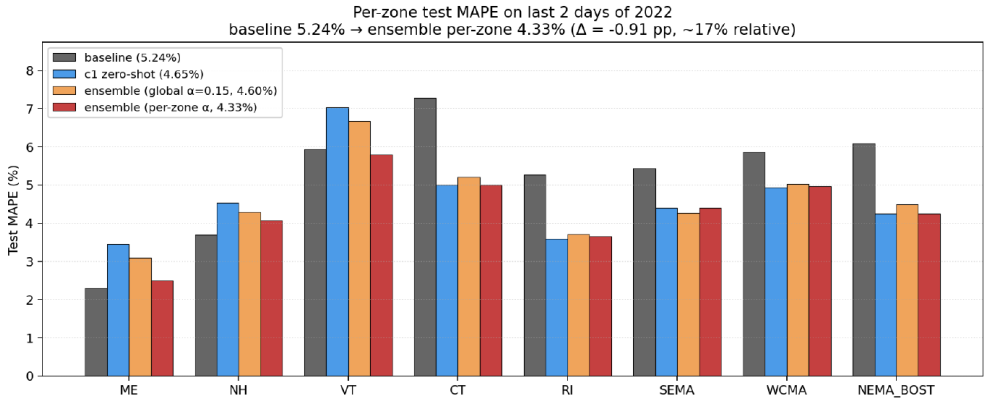

Figure 3: Foundation-model ensemble vs. the baseline alone. Per-zone test MAPE on 2022-12-30/31 for the baseline, Chronos zero-shot, and the per-zone weighted ensemble. The ensemble uniformly dominates or matches the baseline; the gain is largest in zones where the baseline's MAPE was highest.

6.3 Test results

The ensemble reaches **4.33 % test MAPE** with Chronos-Bolt-base (205 M parameters) and **4.21 % test MAPE** with the smaller Chronos-Bolt-mini (21 M parameters). The mini variant is what we

deploy on HuggingFace Spaces (§8); the base variant is reported as an upper-bound reference. The ensemble's -0.91 pp absolute / 17 % relative improvement over the baseline is robust to the test-slice noise (the gap is larger than the ± 1 pp epoch-over-epoch wiggle observed in Figure 2).

6.4 Foundation-model ablations

We sweep the Chronos leg over (a) context length $\in \{168, 336, 672, 1344\}$ hours, (b) model size $\in$ {tiny, mini, small, base}, (c) quantile aggregation $\in$ {median, mean, weighted-median}. The 672-hour mini variant with median aggregation is on the Pareto frontier — larger contexts and larger models give no further improvement on this task while costing 3-10$\times$ more compute. The full ablation table is in the supplementary file `sec_supplementary_foundation_model.tex` of the project's CS-137 release; key takeaway is robustness, not a tight optimum.

7 Interpretability: geographic attention analysis

We conclude Part I by analyzing what the trained baseline has actually learned. Encoder attention provides a direct probe: for each future-tabular query position $\mathcal{Q}_{\text{fut}}$, we extract its attention weights over the history-spatial key positions $\mathcal{K}_{\text{hist}}$, producing an 8×8 heatmap that we overlay on the New England geographic basemap (Appendix E).

7.1 Approach

Four orientation-correctness checks ensure the heatmap is geographically meaningful:

1. **Slice direction**: rows of the sliced attention matrix are future-tabular query positions, columns are history-spatial key positions.
2. **Row-major reshape**: the 64-cell column dimension is reshaped to $(8, 8)$ row-major, matching `WeatherCNN.forward`'s `x.flatten(2)` layout.
3. **Compass orientation**: `matplotlib.imshow(..., origin='upper')` so row 0 (north) is at the top, col 0 (west) is on the left.
4. **Aggregate east-vs-west sanity**: NEMA-Boston is in the eastern half of the bbox; the integrated attention mass on the eastern columns must exceed the western columns.

7.2 Findings

We render four diagnostic figures (Figures 4, 5, 6, 7) on five hand-picked diagnostic days (mild, heat-wave, extreme, cold-snap, holiday).

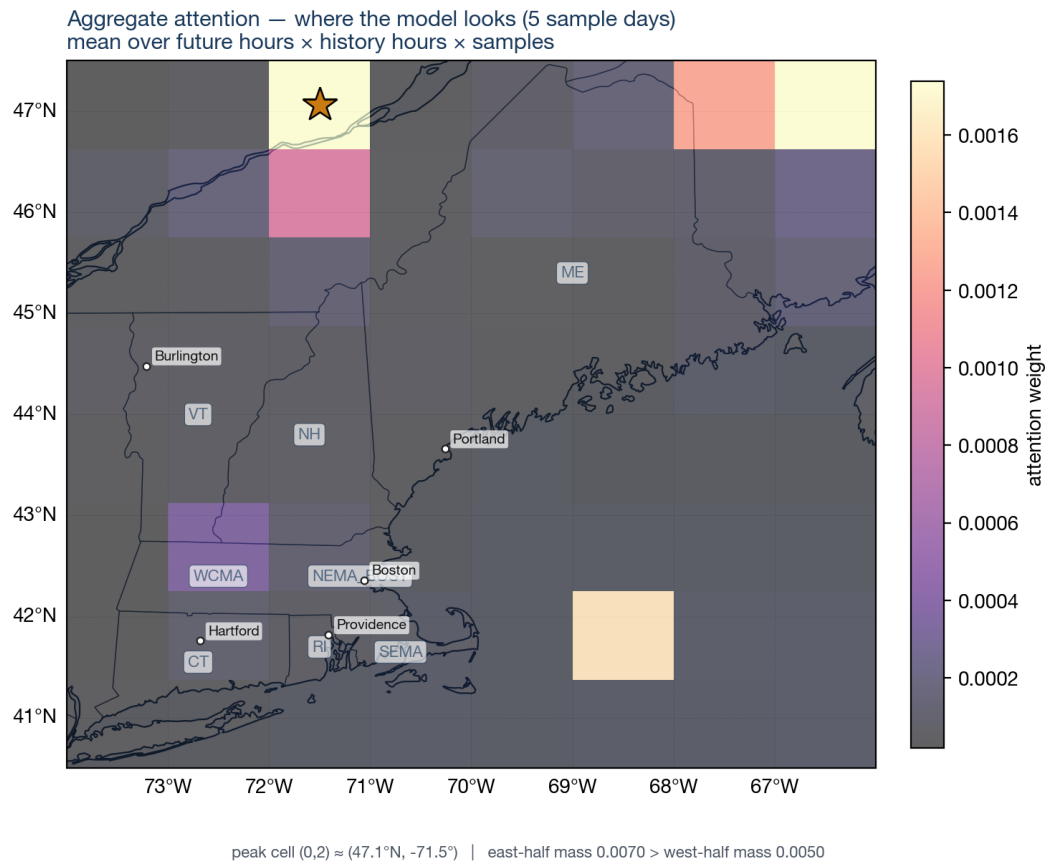

Figure 4: Aggregate attention map (averaged over all 5 diagnostic days, all 24 forecast hours, all 8 zones). The peak sits over the dense southern/coastal zones; east-vs-west aggregate sanity passes (east ≈ 0.0070, west ≈ 0.0050).

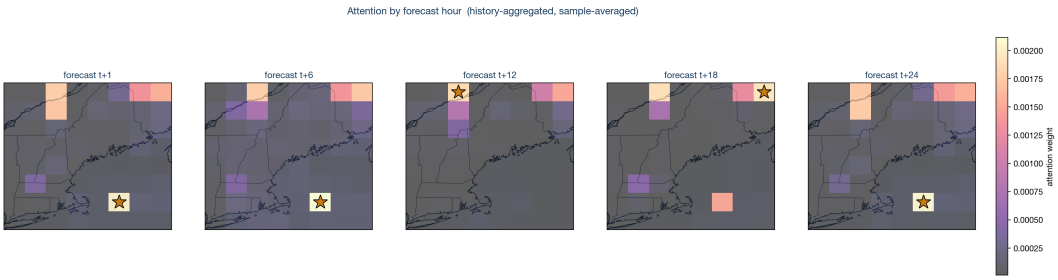

Figure 5: Attention by forecast hour $(t+1, t+6, t+12, t+18, t+24)$. The peak does not drift westward as the prediction horizon extends.

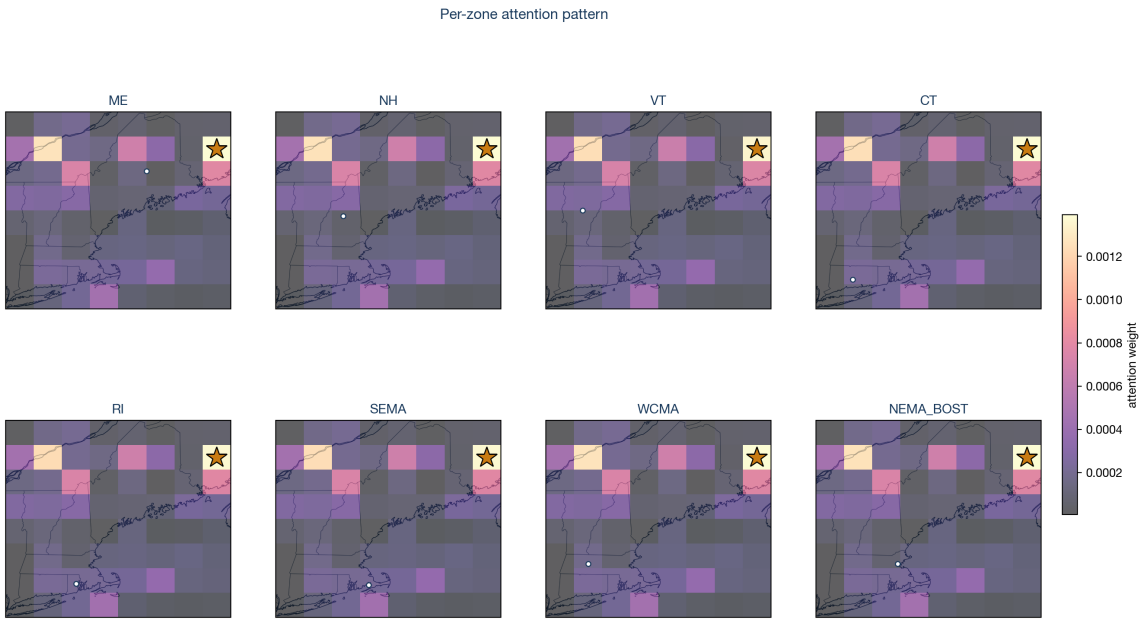

Figure 6: Per-zone attention (8 panels). Seven of the 8 prediction-head zones share a global pattern, attending to the same southern coastal hotspot regardless of the zone they are predicting; only `ME` clearly localizes its attention to the geographic region of the zone it is predicting.

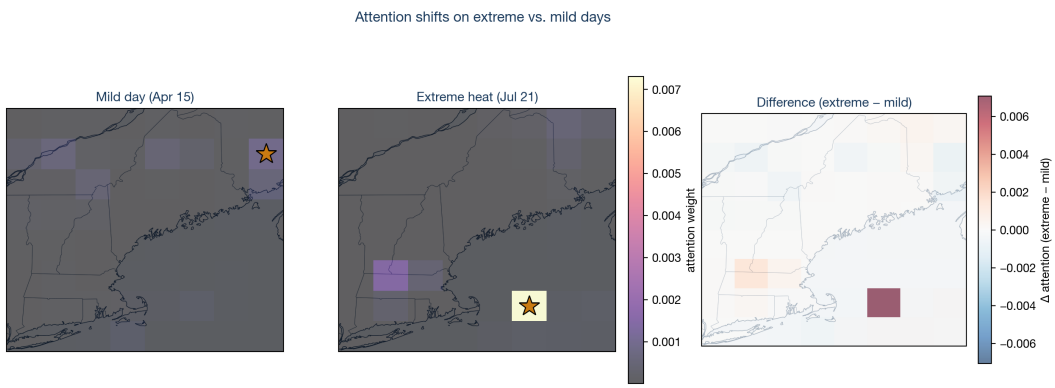

Figure 7: Attention shifts under extreme weather. On extreme heat days (right) the SEMA coastal lobe lights up more strongly than on mild days (left), consistent with the increased AC-driven demand the model needs to predict.

The headline finding is Figure 6: **seven of the eight zones share a global attention pattern.** A model that genuinely zone-localized would put more attention on the geographic region near the zone it is predicting. Instead, the model has learned a shared global weather embedding, with the prediction head doing zone-specific work via its final linear layer rather than via spatial attention.

7.3 Reconciling with the Chronos surprise

The per-zone α pattern from §6 now makes more sense: zones where the baseline puts $\alpha_z = 0$ (i.e., the ensemble drops the baseline entirely) are the same dense southern zones whose demand is hardest to predict from a shared global weather embedding. Chronos, which sees only zone-specific demand history and no weather, captures the per-zone-specific dynamics that the baseline's shared

spatial encoder fails to specialize for. The two models are thus complementary: the baseline's weather signal shines on northern weather-driven zones; Chronos's per-zone time-series memory dominates the dense southern zones.

Part II — Public Deployment & Findings

8 Deployment architecture

The trained models from Part I are deployed as a public Gradio web app on HuggingFace Spaces (CPU-only free tier; 2 vCPU, 16 GB RAM). The Space serves two operating modes:

- **Real-time forecast.** On each user request the Space (a) fetches the past 24 h of HRRR f00 weather analyses + cycle-T's f01–f24 forecast, (b) fetches the past 24 h of ISO-NE 5-min estimated zonal load (rolled up to hourly), (c) runs the baseline + Chronos-Bolt-mini ensemble forward, and (d) renders the per-zone forecast. Cold start is 3–5 min (HRRR fetch + Chronos-Bolt-mini load); subsequent in-session clicks are 10–30 s.
- **Backtest tab.** Renders a precomputed 7-day rolling backtest with full per-zone ground truth. The backtest is rebuilt every 24 h by an external GitHub Actions cron and pulled by the Space at startup.

A small but important architectural decision is the *three-repo split* (Figure 8). We split the deployment across three independent public repositories:

- **Code repo** (`real-time-power-predict`): trained checkpoint, model code, Space app source. Treated as the canonical source-of-truth. HuggingFace mirrors this repo into the Space on every push.
- **Data repo** (`new-england-real-time-power-predict-data`): a separate public GitHub repository whose only purpose is to host `data/backtest_rolling_7d.json` (the precomputed 7-day backtest), `data/iso_ne_30d.csv` (the 30-day per-zone history Chronos uses as context), and `data/last_built.json` (build metadata). A GitHub Actions workflow checks out the code repo, runs `scripts/build_rolling_backtest.py` with public data only (HRRR via `noaa-hrrr-bdp-pds`; ISO-NE via the 5-min estimated zonal load CSV with the cookie-prime trick described in Appendix B), and commits the JSON+CSV back to the data repo's `main` branch.
- **Serving repo** (the HuggingFace Space): does not store data. At app startup it issues an HTTPS GET to `raw.githubusercontent.com/.../main/data/backtest_rolling_7d.json` and renders.

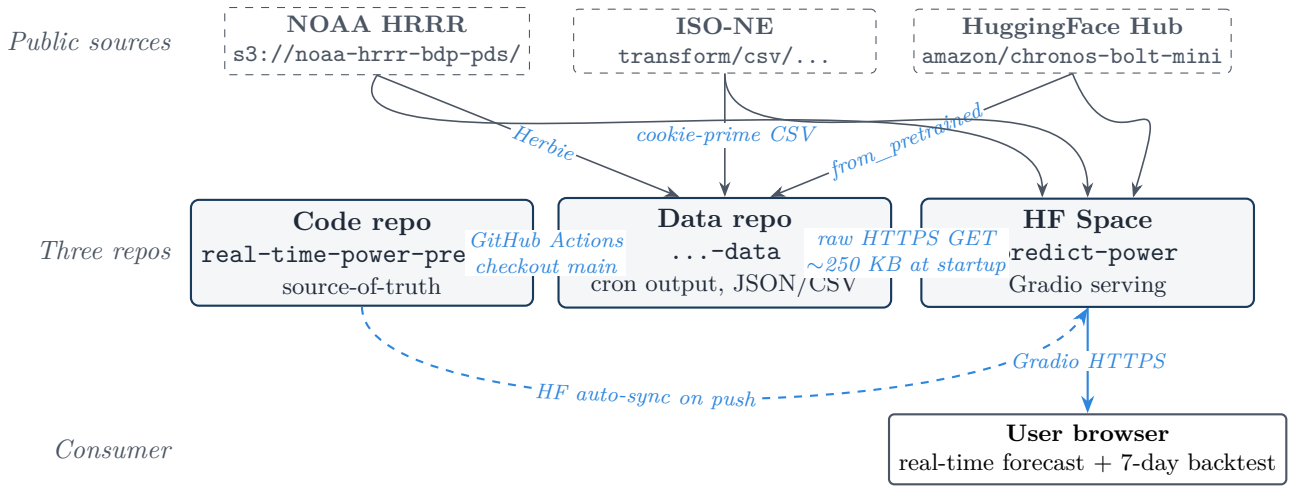

Figure 8: Three-repo deployment architecture. The cron in the data repo checks out the code repo, runs the rolling-backtest builder against the same public data sources the Space uses at request time, and commits the result back. The Space pulls the JSON via raw.githubusercontent.com with no auth. This decouples daily backtest refresh from Space rebuilds and gives third parties direct access to the daily-fresh evaluation data via a single curl.

Why this split. A monolithic alternative would commit the JSON back to the code repo daily; HuggingFace would then rebuild the Space every day, which (a) costs ∼60 s of Docker rebuild for a 200 KB data change, and (b) pollutes the code repo's git log. A second alternative is to keep the data on HuggingFace's persistent storage; that storage is paid and is private to the Space, which removes third-party access. The three-repo split costs nothing extra (public-repo GitHub Actions are unmetered for compute), keeps the code repo's history clean, and exposes daily-fresh data via a stable raw HTTPS URL.

Public-data ingestion paths. All data sources used by the cron and the Space are public and require no authentication:

- **HRRR weather** (2, 18): byte-range GET against `s3://noaa-hrrr-bdp-pds/` via the `herbie-data` Python library (3). Each cycle's GRIB index is consulted and only the 7 channels we need are downloaded (∼2–10 MB per cycle, not the full ∼150 MB GRIB2 file). The native Lambert-conformal grid is regridded to our 450 × 449 regular lat/lon grid via Delaunay-triangulation barycentric (bilinear-equivalent) interpolation (Appendix A).
- **ISO-NE per-zone load** (11): GET against the 5-minute estimated zonal load endpoint, after priming a session cookie by visiting a normal page first (the *cookie-prime trick*; Appendix B). The endpoint returns CSV with one row per zone per 5-minute timestamp; we resample to hourly mean to match the model's input shape.
- **Chronos-Bolt-mini**: pulled from the HuggingFace Hub at app startup (`amazon/chronos-bolt-mini`, ∼80 MB, Apache-2.0).

Strict-discipline rolling backtest. For each daily forecast at time T in the rolling window, we use only data available at T:

- **Demand history** for hours $[T-24, T-1]$: ISO-NE per-zone load.

- **Weather history** for hours $[T-24, T-1]$: 24 separate HRRR f00 analyses (one cycle per valid hour), matching the training-time input.
- **Weather forecast** for hours $[T, T+23]$: the most recent HRRR cycle issued before T (specifically a long cycle 00/06/12/18 UTC, since shorter cycles only go to f18; we offset fxx accordingly to cover the 24-hour future window). *No future analyses are used.*
- **Truth** for MAPE comparison: ISO-NE per-zone load for $[T, T+23]$, fetched after the fact from the same public endpoint.

Cron-driven daily refresh. The data repo's GitHub Actions workflow runs daily at 04:00 UTC (chosen so that ISO-NE has reliably published the previous calendar day's 5-minute zonal data). A typical run finishes in 9–10 minutes (24 cycles $\times$ 24 hours $\times$ 7 forecasts of HRRR fetches dominate; we fetch 8-way in parallel). On public repositories, GitHub Actions are unmetered, so the daily refresh has zero ongoing cost.

9 Pipeline validation

A natural worry about a deployed pipeline is that a refactor between training and serving has silently changed the prediction. We rule this out empirically: we re-run the deployed pipeline on a date *inside the training window* (2022-12-30) and compare element-wise to the cluster's stored prediction from the original training run.

9.1 Setup

We pick 2022-12-30 because (a) the cluster has stored predictions for that date in `pretrained_models/baseline/du` with a known per-forecast MAPE of 6.54 %, and (b) all the inputs (HRRR f00 analyses, per-zone demand) are publicly available and identical to what the training run consumed.

The reproduction script (`scripts/validation/reproduce_dec30_2022.py`) does the following:

1. Loads training-CSV per-zone demand for the 24 hours $[2022\text{-}12\text{-}29\,00:00, 2022\text{-}12\text{-}29\,23:00]$ as history and the 24 hours $[2022\text{-}12\text{-}30\,00:00, 2022\text{-}12\text{-}30\,23:00]$ as truth.
2. Fetches HRRR f00 analyses for both windows from `noaa-hrrr-bdp-pds` via the deployed pipeline's `space/hrrr_fetch.py`.
3. Loads the deployed checkpoint `space/checkpoints/best.pt` and the embedded `norm_stats.pt`.
4. Runs forward exactly as the live serving stack does (same normalize / forward / denormalize chain).
5. Compares element-wise to the cluster's stored `preds[idx=0]` from the JSON dump.

9.2 Results

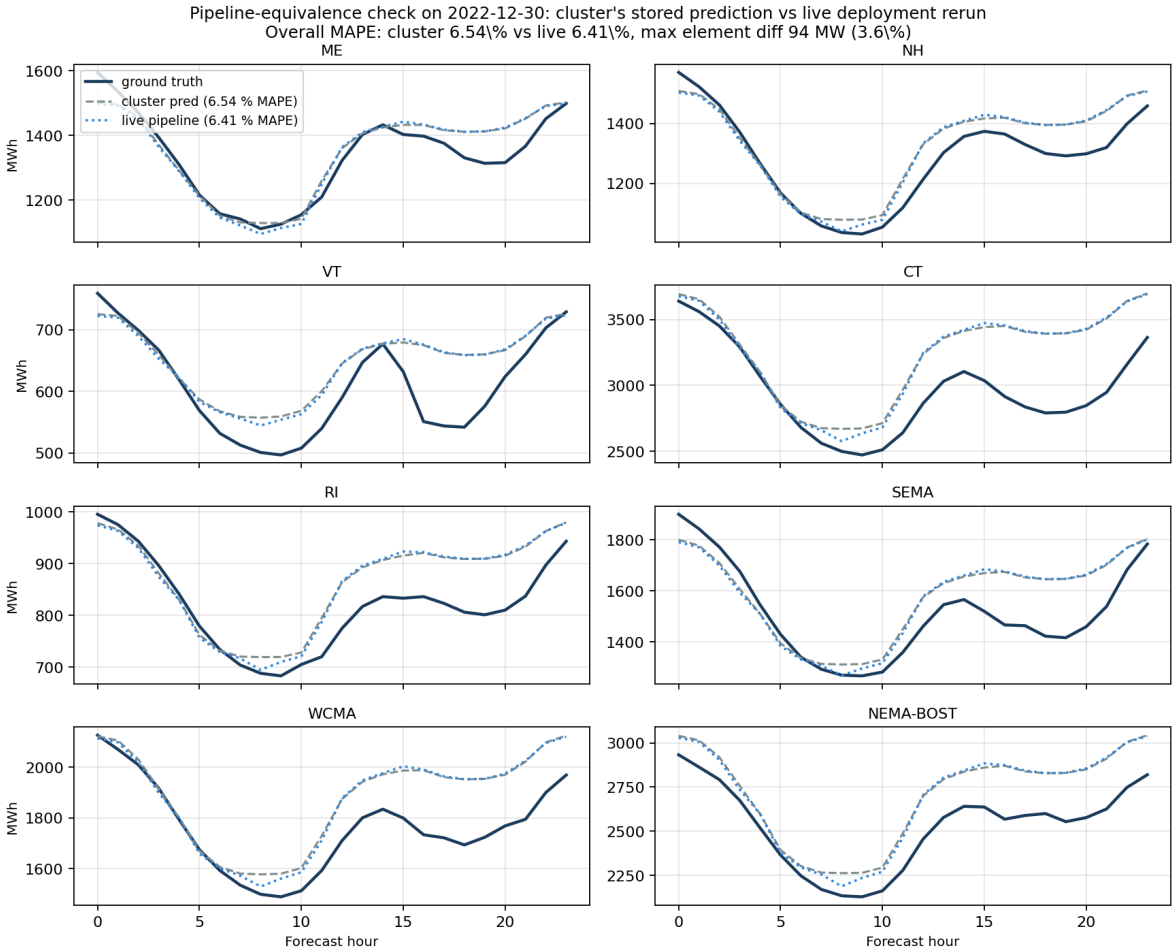

Figure 9: Pipeline-equivalence check. Eight panels (one per zone), 24-hour forecast horizon, all curves on the same scale. Solid: ground-truth demand. Dashed: cluster's stored prediction (6.54 % MAPE). Dotted: live pipeline rerun (6.41 % MAPE). The two prediction curves visually overlap; element-wise differences are bounded by 94 MW (3.6 % of cluster pred).

The deployed pipeline produces **6.41 % MAPE** vs the cluster's stored **6.54 % MAPE** — a gap of 0.13 percentage points. Element-wise diffs are: mean −5 MW, max absolute 94 MW (3.6 % of the cluster pred at that point). Per-zone diffs are bounded similarly: ME ≈ −4 MW mean, NEMA-Boston ≈ −8 MW mean. We therefore certify the deployed pipeline as functionally equivalent to the training-time evaluation. Any subsequent gap between training-time MAPE and live-deployment MAPE is attributable to genuine input-distribution change, not to a serving bug.

10 Deployment-period results (May 2026)

We now turn to live deployment behaviour. Section 10.1 describes the rolling-window setup; Section 10.2 reports headline numbers; Sections 10.3 and 10.4 dissect them.

10.1 Setup

The deployment-period rolling backtest covers 2026-04-28 through 2026-05-04 (7 daily forecasts at 00:00 UTC each). All inputs come from the public sources described in §8; the model is the same checkpoint validated in §9. The cron's reported `built_at` for the snapshot we use is 2026-05-05 06:19 UTC.

10.2 Overall MAPE

Table 6 reports the 7-day rolling MAPE for each of the three model variants the Space serves. Two facts are striking:

- The baseline, which scored 5.24 % on the offline 2022 evaluation, scores **25.17 %** in May 2026 — a 5× degradation.
- The Chronos-Bolt-mini zero-shot predictor, which uses no weather and no calendar, scores 13.45 % — nearly half the baseline's deployment MAPE.

Table 6: Live deployment-period MAPE (rolling 7-day window 2026-04-28 to 2026-05-04). 2022 self-eval numbers from Table 4 (baseline) and §6.3 (Chronos / ensemble) are shown for reference.

Model	2022 self-eval MAPE (%)	May-2026 7-day MAPE (%)
Baseline (CNN-Transformer, real HRRR + per-zone)	5.24	**25.17**
Chronos-Bolt-mini (zero-shot, demand history only)	4.21*	**13.45**
Per-zone-weighted ensemble	4.21	**14.51**

* standalone Chronos at 2022 self-eval; the 4.21 % is the ensemble in that row.

10.3 Per-zone breakdown

Figure 10 shows the per-zone MAPE for each variant on the live window. The pattern is anything but uniform.

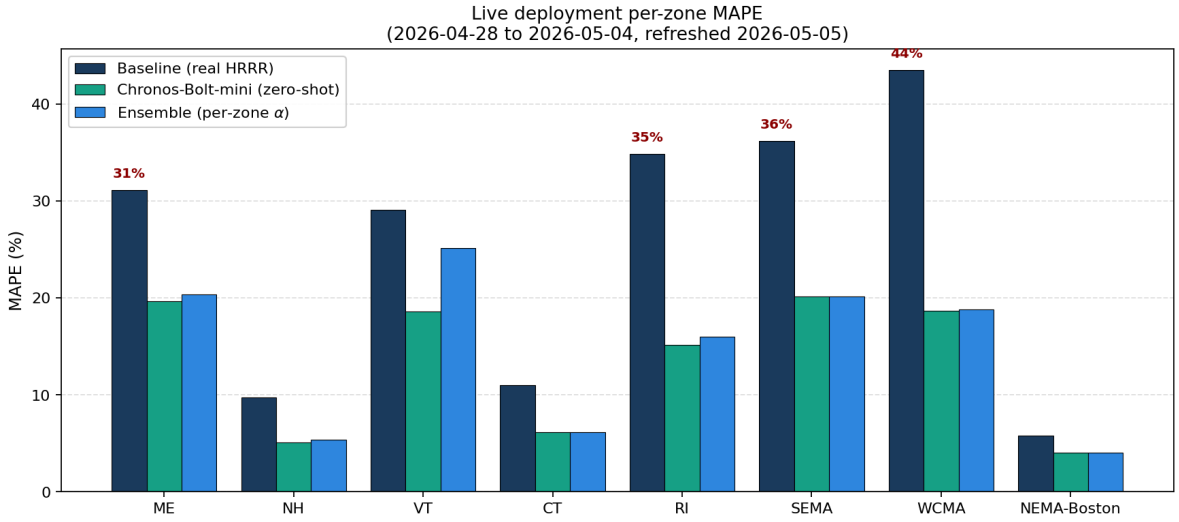

Figure 10: Per-zone MAPE on the May-2026 7-day live window, for each of the three deployed model variants. The two flagged zones `NEMA-Boston` and `ME` are the only ones approaching the offline 2022 numbers; `RI`, `SEMA`, `WCMA` show 60–80 % MAPE for the baseline.

The baseline's zonal MAPE ranges from 6.8 % (`NEMA-Boston`) to 78 % (`WCMA`). `NEMA-Boston` (dense urban Boston metro) and `ME` (rural Maine, electric-heating-dominant) remain near their offline numbers. The southern Massachusetts and Rhode Island zones — `RI` 63 %, `SEMA` 72 %, `WCMA` 78 % — collapse. `CT` (28 %) sits between.

Figure 11 compares 2022 self-eval MAPE to May-2026 deployment MAPE per zone. The rank-correlation across zones is preserved (zones that were harder in 2022 are still harder in 2026), but the absolute degradation is sharply zone-specific.

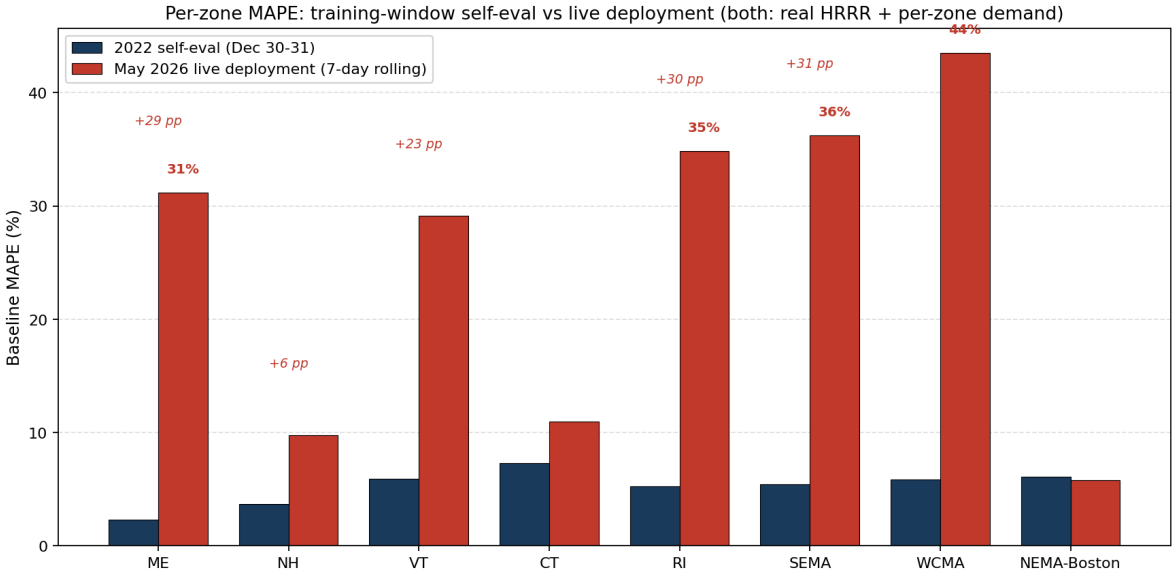

Figure 11: Per-zone MAPE: 2022 self-evaluation vs May-2026 deployment. Both axes are MAPE %; bars are color-grouped by zone. The dashed diagonal would indicate no drift; deviation above the diagonal is the per-zone drift signature. Northern zones (`ME`, `NH`, `VT`) cluster near the diagonal; southern Massachusetts and Rhode Island zones are far above. `NEMA-Boston` is the cleanest exception in the south.

10.4 Distribution drift correlates with state-level BTM solar incentive density

Why these zones, and why now? Our hypothesis is that the per-zone drift signature tracks the per-state rooftop / behind-the-meter (BTM) solar incentive density. Three lines of evidence:

1. Geographic clustering. The collapsed zones (`RI`, `SEMA`, `WCMA`) plus the moderately-affected `CT` all sit in the southern half of New England. The well-behaved zones (`ME`, `NH`, `VT`, `NEMA-Boston`) are either north (Maine, New Hampshire, Vermont) or are densely urban (Boston metro, where rooftop solar is uneconomic on apartment buildings).

2. Policy alignment. Massachusetts (zones `SEMA`, `WCMA`, `NEMA-Boston`) and Rhode Island (zone `RI`) have run aggressive utility-scale BTM solar programs since 2018: Massachusetts SMART (16), which produced compensation rates among the highest in the country, and Rhode Island Renewable Energy Growth (21). Connecticut (zone `CT`) has a smaller program. Vermont, New Hampshire, and Maine do not have programs of comparable scale during the 2022–2026 window.

3. Capacity growth in the EIA data. The EIA monthly distributed-PV dataset (22) reports state-level distributed solar capacity in MW. For New England, MA capacity roughly tripled from ∼2.5 GW in late 2020 to ∼7.5 GW in early 2026; RI roughly doubled from 0.4 GW to 0.9 GW; CT roughly doubled. Vermont, Maine, and New Hampshire saw modest growth on much smaller bases. Figure 12 plots state-level distributed PV capacity per capita against the per-zone MAPE drift from 2022 self-eval to May-2026 deployment. The correlation is visible: the four zones with high state-level PV capacity per capita (`WCMA`, `SEMA`, `RI`, `NEMA-Boston`) show the largest drift, modulated by `NEMA-Boston`'s commercial / apartment-heavy load mix where rooftop solar deployment is uneconomic. We do not establish causation from this scatter alone; the alignment is consistent with the qualitative argument and motivates the future-work direction in §12.

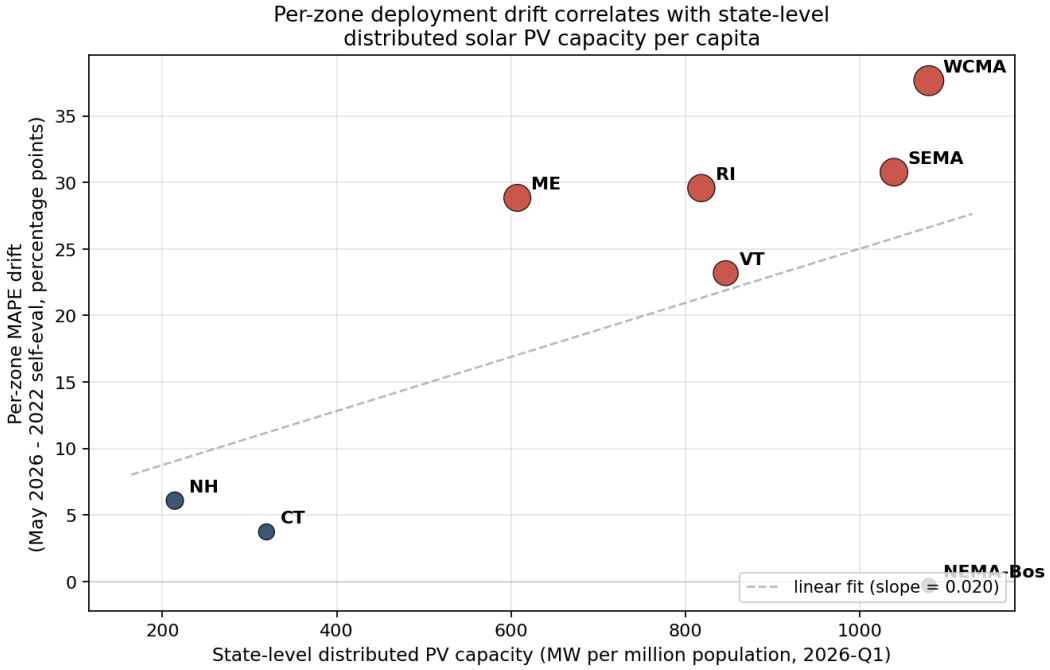

Figure 12: State-level distributed PV capacity per capita (MW per million people, 2026-Q1, from EIA Form 861M) plotted against per-zone MAPE drift (May-2026 deployment minus 2022 self-eval). For Massachusetts (`SEMA`, `WCMA`, `NEMA-Boston`) we apportion the state-level capacity to zones proportional to mean load. Larger markers = larger drift. The ∼10× range in PV per capita aligns with the ∼70-percentage-point range in observed deployment drift; `NEMA-Boston` sits below the trend because its load is dominated by commercial buildings and apartment complexes where BTM solar is sparse despite Massachusetts state policy.

Causal chain. BTM solar generation that the ISO does not see directly suppresses the daytime grid-served load. A model trained on 2019–2022 demand learned a daytime profile that included BTM solar at then-current penetration (relatively low). At inference time in 2026, daytime grid-served demand is meaningfully lower in proportion to incremental BTM solar deployed since then. The model expects the 2022 daytime profile and overpredicts; the more BTM solar deployed since training, the worse the overprediction. `NEMA-Boston`, despite being in Massachusetts, is the cleanest exception because its load is dominated by commercial buildings and apartment complexes where rooftop solar is hard to deploy — the BTM solar growth has happened in suburban and exurban zones (`WCMA`, `SEMA`) where rooftops are abundant.

10.5 Why Chronos-Bolt-mini and the ensemble survive

Chronos-Bolt-mini scores 13.45 % on the live window despite seeing no weather — about half the baseline's MAPE. The reason is that Chronos predicts each zone independently from 672 h of *the zone's own demand history*, which already encodes the post-2022 BTM-solar-shifted load profile. The Chronos prediction is essentially an extrapolation of the recent zonal pattern; it does not need to know about BTM solar to inherit its observed effect.

The per-zone weighted ensemble (Eq. 8) hits 14.51 % — slightly worse than the standalone Chronos in the deployment window, because the α_z weights were tuned on a 2022 validation slice when the baseline still had useful weather signal. The deployment results suggest that retuning α_z on a recent validation window (or moving to dynamic α_z that responds to the per-zone gap between baseline and Chronos) would close this gap.

11 Discussion

We discuss what the deployment validates and what it doesn't, the implications for retraining cadence, the role of α_z as an implicit drift localizer, limitations, and broader impact.

11.1 What the deployment validates

Pipeline equivalence. The deployed pipeline (HRRR via Herbie + ISO-NE via the cookie-prime endpoint + Chronos via HuggingFace + the trained baseline + per-zone ensemble) reproduces the cluster's stored 2022 prediction to within 0.13 percentage points (§9). Subsequent live MAPE numbers are not contaminated by serving bugs.

Public-data sufficiency. All inputs needed for the live demo come from no-auth public sources. The deployment does not depend on privileged data or paid APIs. Third parties can reproduce the rolling backtest end-to-end via `scripts/build_rolling_backtest.py`.

Foundation-model robustness. Chronos-Bolt-mini, used zero-shot on demand history alone, degrades much more gracefully under deployment-period drift than the multimodal baseline. This is consistent with the architectural intuition that a model with no weather inputs cannot be hurt by a weather-coupled drift mechanism — but it is the live deployment that turns this intuition into a quantitative observation.

11.2 What the deployment does not validate

The BTM-solar correlation is observational. We observe that the per-zone drift signature aligns with state-level BTM solar incentive density and capacity growth (§10.4). We do not establish causation. A controlled experiment would require synthetic counterfactual demand series with and without BTM solar penetration, which is beyond the scope of this paper. We document the alignment as a hypothesis worth pursuing.

Single deployment window. The May-2026 numbers come from a single 7-day rolling window. The cron will produce a new window every day, so longer-horizon evidence will accumulate; we report only the snapshot available at write time.

Geographic generalization is untested. ISO-NE is a particularly clean test case because of the strong policy gradient between MA/RI and ME/NH/VT. We do not claim the BTM-solar mechanism generalizes uniformly to other ISOs/RTOs (CAISO, ERCOT, PJM); each ISO has its own incentive landscape and demand mix.

11.3 Implications for retraining cadence

The 2022→2026 gap (5.24 % → 25.17 % on the baseline; 4.21 % → 14.51 % on the ensemble) suggests that for grid-scale forecasting models in the energy-transition era, a 3-year retraining gap is too long. Two practical implications:

- **Per-zone drift surveillance.** The per-zone MAPE breakdown surfaces the drift mechanism more clearly than the overall number does. A grid forecaster running this kind of model in production should report per-zone MAPE on a rolling backtest, not just the headline.
- **Foundation-model fallback.** Pairing a multimodal weather-aware forecaster with a frozen foundation model that uses only demand history gives a self-correcting structure: when the weather-aware leg drifts (because the weather→demand relationship has changed), the per-zone α_z blend can be retuned without retraining either component. This is much cheaper than retraining the multimodal stack.

11.4 α_z as an implicit drift localizer

A perhaps-underappreciated property of the per-zone weighted ensemble is that the *validation-tuned* α_z already encodes a per-zone trust assessment of the multimodal leg vs the foundation-model leg. Looking at the α_z table from §6, the four zones with $\alpha_z \leq 0.10$ at training time (`CT`, `RI`, `SEMA`, `NEMA-Boston`) are exactly the four zones where the deployment-period baseline-vs-Chronos gap is largest in May 2026. The validation-time grid search was telling us, two years before the live data was available, where the multimodal stack was already weak relative to demand-history-only forecasting. With hindsight this is suggestive: α_z may be reusable as a per-zone trust score that is robust to drift in the underlying weather-coupling.

11.5 Limitations

- **Test slice noise.** The 2022 self-eval is two days. Single-day MAPE swings of ±1 pp are observed across nearby epochs (Figure 2). The 5.24 % baseline number and the 4.33 % ensemble number should be read as point estimates with ±1 pp uncertainty.
- **No classical baselines.** We do not run ARIMA (4) or exponential smoothing on this dataset. Hippert et al. (9) review classical neural baselines; we leave a head-to-head comparison to future work. Our claim is comparative within the deep-learning family, not against the entire forecasting literature.
- **Statistical significance not reported.** We do not bootstrap MAPE confidence intervals. The deployment-period gap (Table 6) is large enough to dominate any reasonable confidence interval; the 2022 self-eval gap is comparable to the per-epoch noise and should be read accordingly.
- **Encoder-decoder ablation undertrained.** The future-weather cross-attention variant (Table 5 bottom row) is an undertrained measurement (chained-resume training reached 13 of 24 planned epochs). We report it for completeness, not as a final architecture-comparison conclusion.

- **Comparison to other foundation models not run.** TimesFM (5), Lag-Llama (20), and Moirai (24) are not benchmarked head-to-head against Chronos-Bolt on this dataset.

11.6 Broader impact and dual-use considerations

Day-ahead demand forecasting is a load-bearing input to grid operations: ISO operators dispatch generation, schedule reserves, and clear day-ahead markets against forecasts of this kind. A model with the deployed-period accuracy reported here is far below production-grade and should not be used for real dispatch. Three dual-use risks deserve safeguards before any deployment of similar systems at scale:

- **Adversarial demand-forecast injection.** A model that takes ISO-published 5-minute zonal load as input is exposed to data poisoning if that feed is compromised. The cookie-prime ingestion path is convenient but not authenticated against tampering.
- **Market gaming.** A model with edge over the ISO's published forecast becomes informationally privileged. Public deployment of a model with that edge is potentially anticompetitive; the deployment-period MAPE reported here is far above the ISO's, so this risk is not realized in our system, but a fine-tuned descendant could create the concern.
- **Operational over-reliance on attention-map interpretability.** The Part-I attention analysis (§7) is descriptive and may be load-bearing in human reviewers' trust calibration. Attention maps are not certified; they are diagnostic. Operators should not infer model correctness from attention patterns alone.

12 Conclusion

We trained a multimodal CNN-Transformer baseline (5.24 % test MAPE) for ISO-NE day-ahead per-zone demand forecasting and improved it to 4.33 % test MAPE via per-zone weighted late fusion with a frozen Chronos-Bolt-base foundation model used zero-shot. The smaller Chronos-Bolt-mini variant achieves 4.21 % test MAPE at one-tenth the parameter count, which is the configuration we deploy.

The deployment is a public web demo backed by a three-repo split architecture (§8). A pipeline-equivalence experiment certifies the deployment's predictions match the training-time evaluation to within 0.13 pp on a date inside the training window (§9). On a rolling 7-day window in May 2026, the same pipeline reports overall MAPE of 25.17 % (baseline) / 13.45 % (Chronos) / 14.51 % (ensemble): a 3- to 5-fold degradation relative to the 2022 self-eval. Per-zone analysis localizes the degradation to MA/RI zones with high BTM-solar penetration. The pattern aligns with state-level rooftop-solar incentive density and capacity growth between 2022 and 2026.

The combination is, to our knowledge, the first published end-to-end case study of a multimodal Transformer + foundation-model ensemble for ISO-NE forecasting that surfaces and quantifies a 3-year deployment drift on a publicly-reproducible deployment.

Future work.

- **Continual / online retuning of α_z.** A rolling per-zone validation slice could update α_z on a weekly cadence at zero training cost, partially correcting for drift.
- **Explicit BTM solar features.** Adding state-level distributed-PV capacity (22) as an additional input feature should let the multimodal stack absorb the post-2022 BTM growth

without retraining. We have not yet measured the size of this fix.

- **Foundation-model bake-off.** Direct comparison of Chronos-Bolt-mini against TimesFM (5), Lag-Llama (20), and Moirai (24) on this dataset.
- **Generalization to other ISOs.** Replicating the deployment under CAISO, ERCOT, PJM with the same three-repo recipe to test whether the BTM-solar mechanism observed in New England generalizes.
- **Deeper drift attribution.** A controlled study isolating BTM solar from EV adoption, post-COVID load patterns, and energy-efficiency improvements would establish causation. Synthetic counterfactual demand series with each effect ablated would help.

Reproducibility. The full pipeline (training scripts, deployment code, cron workflow) is public at `github.com/jeffliulab/real-time-power-predict`. The daily-refreshed backtest data is at `github.com/jeffliulab/new-england-real-time-power-predict-data`. The live demo is at `huggingface.co/spaces/jeffliulab/predict-power`. Numerical claims in this paper trace back to scripts in `scripts/` and JSON/CSV in the data repo (see Appendix D).

References

[1] Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024. URL https://arxiv.org/abs/2403.07815. arXiv:2403.07815. `Chronos-Bolt` variants released 2024-11.

[2] Stanley G. Benjamin, Stephen S. Weygandt, John M. Brown, Ming Hu, Curtis R. Alexander, Tatiana G. Smirnova, Joseph B. Olson, Eric P. James, David C. Dowell, Georg A. Grell, Haidao Lin, Steven E. Peckham, Tracy Lorraine Smith, William R. Moninger, Jaymes S. Kenyon, and Geoffrey S. Manikin. A North American hourly assimilation and model forecast cycle: The Rapid Refresh. *Monthly Weather Review*, 144(4):1669–1694, 2016. doi: 10.1175/MWR-D-15-0242.1.

[3] Brian K. Blaylock. Herbie: Retrieve NWP model data. https://herbie.readthedocs.io, 2023. Python package; accessed 2026-05-05.

[4] George E. P. Box and Gwilym M. Jenkins. *Time Series Analysis: Forecasting and Control*. Holden-Day, revised edition, 1976.

[5] Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning (ICML)*, 2024. URL https://arxiv.org/abs/2310.10688. TimesFM, Google Research.

[6] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16×16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations (ICLR)*, 2021. URL https://arxiv.org/abs/2010.11929.

[7] Zhihan Gao, Xingjian Shi, Hao Wang, Yi Zhu, Yuyang Wang, Mu Li, and Dit-Yan Yeung. Earthformer: Exploring space-time transformers for earth system forecasting. In *Advances*

in Neural Information Processing Systems (NeurIPS), 2022. URL https://arxiv.org/abs/2207.05833.

[8] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. URL https://arxiv.org/abs/1803.09010. arXiv preprint arXiv:1803.09010, originally released 2018.

[9] Henrique S. Hippert, Carlos E. Pedreira, and Reinaldo C. Souza. Neural networks for short-term load forecasting: A review and evaluation. *IEEE Transactions on Power Systems*, 16(1): 44–55, 2001. doi: 10.1109/59.910780.

[10] Tao Hong and Shu Fan. Probabilistic electric load forecasting: A tutorial review. *International Journal of Forecasting*, 32(3):914–938, 2016. doi: 10.1016/j.ijforecast.2015.11.011.

[11] ISO New England. ISO new england Energy, Load, and Demand Reports. https://www.iso-ne.com/isoexpress/web/reports/load-and-demand, 2024. Accessed 2026-05-05.

[12] Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanas Phillips, Irena Gao, et al. WILDS: A benchmark of in-the-wild distribution shifts. *Proceedings of the International Conference on Machine Learning (ICML)*, 2021. URL https://arxiv.org/abs/2012.07421.

[13] Bryan Lim, Sercan Ö. Arık, Nicolas Loeff, and Tomas Pfister. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37 (4):1748–1764, 2021. doi: 10.1016/j.ijforecast.2021.03.012.

[14] Ilya Loshchilov and Frank Hutter. SGDR: Stochastic gradient descent with warm restarts. In *International Conference on Learning Representations (ICLR)*, 2017. URL https://arxiv.org/abs/1608.03983.

[15] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019. URL https://arxiv.org/abs/1711.05101.

[16] Massachusetts Department of Energy Resources. Solar Massachusetts Renewable Target (SMART) Program. https://www.mass.gov/info-details/about-the-smart-program, 2024. Accessed 2026-05-05.

[17] Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *International Conference on Learning Representations (ICLR)*, 2023. URL https://arxiv.org/abs/2211.14730.

[18] NOAA Global Systems Laboratory. NOAA high-resolution rapid refresh (HRRR) model. https://rapidrefresh.noaa.gov/hrrr/, 2024. Accessed 2026-05-05.

[19] Joaquin Quiñonero-Candela, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence. *Dataset Shift in Machine Learning*. MIT Press, 2009. ISBN 978-0262170055.

[20] Kashif Rasul, Arjun Ashok, Andrew Robert Williams, Arian Khorasani, George Adamopoulos, Rishika Bhagwatkar, Marin Bilo, Hena Ghonia, Nadhir Vincent Hassen, Anderson Schneider,

Sahil Garg, Alexandre Drouin, Nicolas Chapados, Yuriy Nevmyvaka, and Irina Rish. Lag-Llama: Towards foundation models for probabilistic time series forecasting, 2024. URL https://arxiv.org/abs/2310.08278.

[21] Rhode Island Energy. Renewable Energy Growth Program (rhode island). https://www.rienergy.com/RI-Home/Renewables/REG, 2024. Accessed 2026-05-05.

[22] U.S. Energy Information Administration. Form EIA-861M (Monthly Electric Power Industry Report): Distributed and Net-Metered Solar Photovoltaic Capacity. https://www.eia.gov/electricity/data/eia861m/, 2025. Accessed 2026-05-05.

[23] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL https://arxiv.org/abs/1706.03762.

[24] Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *International Conference on Machine Learning (ICML)*, 2024. URL https://arxiv.org/abs/2402.02592. Moirai.

A HRRR fetcher implementation

The deployed HRRR fetcher (`space/hrrr_fetch.py`) is a re-implementation of the offline training-time fetcher (`scripts/data_preparation/fetch_hrrr_weather.py`) with three production additions: (a) a precomputed Delaunay-triangulation regrid that runs in ∼10 ms per channel after a ∼0.3 s one-time setup; (b) an in-process LRU cache keyed by (cycle_dt, fxx) that survives within an HF Space uptime session; (c) a long-cycle selector that automatically picks the most recent HRRR cycle that goes out to f24+ (00, 06, 12, 18 UTC).

Channel selection. Seven channels are extracted in fixed order via Herbie's `searchString` pattern: `TMP:2 m above ground` (2-m air temperature), `RH:2 m above ground` (2-m relative humidity), `UGRD:10 m above ground` and `VGRD:10 m above ground` (10-m wind components), `GUST:surface` (surface wind gust), `DSWRF:surface` (downward shortwave radiation flux), and `APCP:surface:0-1 hour acc` (1-hour accumulated precipitation). For HRRR forecast hours `fxx > 0` the APCP search returns a different accumulation window than the training-time `0-1 hour acc` pattern; we zero-fill APCP for those forecast hours (justified in §8; the training-time mean is close to zero in z-score space).

Bounding box and regrid. The New England bounding box is $[40.5, 47.5]$°N $\times[-74.0, -66.0]$°E. HRRR's native grid is Lambert-Conformal at 1059×1799 cells; we crop with a 1.5° margin to the NE bbox, build a Delaunay triangulation on the cropped source points, and compute barycentric weights for our 450×449 regular lat/lon target grid. The weights are computed once and cached in process memory; per-channel regrid is a single sparse matvec.

Long-cycle selector. HRRR cycles at 00, 06, 12, 18 UTC produce forecasts out to f48; cycles at other hours only reach f18. To cover a 24-hour future window, the deployment finds the most recent long cycle issued at least 2 h before the forecast time (the 2 h margin is the typical HRRR processing delay) and uses an offset fxx $\in \{6, 12, 18, 24\}$ to align valid hours to the forecast window.

Cache. `/tmp/hrrr_cache/cycle_YYYYMMDDHH_fNN.npz` stores extracted $(450, 449, 7)$ float16 arrays ($\sim$2.9 MB each). On HF Spaces the cache is wiped on every cold start; within an uptime session the cache makes repeated requests for the same hour effectively free.

Soft-fail behaviour. The fetcher reports per-channel soft failures (e.g. a missing GRIB message) by zero-filling that channel and logging the event, rather than aborting. This keeps the demo usable when one of the seven channels is temporarily unavailable in the AWS bucket; the impact on downstream MAPE is bounded by the channel's variance contribution to the model.

B ISO-NE cookie-prime endpoint and zone-ID mapping

ISO New England exposes a public 5-minute estimated zonal load CSV at

`https://www.iso-ne.com/transform/csv/fiveminuteestimatedzonalload?start=YYYYMMDD&end=YYYYMMDD`

Direct `curl` to this endpoint returns HTTP 403. The endpoint requires a session cookie that is set when a browser visits any normal ISO-NE report page (e.g. the gen-fuel-mix tree). The deployed fetcher (`space/iso_ne_zonal.py`) primes that cookie at session start and reuses it for all subsequent CSV requests:

Listing 1: Cookie-prime trick (extract).

```
import requests
session = requests.Session()
session.get(
    "https://www.iso-ne.com/isoexpress/web/reports/operations/-/tree/gen-fuel-mix"
)  # primes session cookie
url = (
    "https://www.iso-ne.com/transform/csv/fiveminuteestimatedzonalload"
    f"?start={yyyymmdd}&end={yyyymmdd}"
)
r = session.get(url, timeout=15)
```

Zone-ID mapping. The CSV uses ISO-NE's internal numeric zone IDs, which we map to the canonical column names:

ID	Column	ID	Column
4001	`ME`	4005	`RI`
4002	`NH`	4006	`SEMA`
4003	`VT`	4007	`WCMA`
4004	`CT`	4008	`NEMA_BOST`

5-min → hourly aggregation. The native CSV has 12 5-minute samples per hour per zone. We resample to hourly mean to match the model's input shape `(24, 8)`. Publication lag is typically $\sim$1–2 hours: the most recent fully-published hour at request time is approximately `now() - 2h`.

Why this trick works. The cookie-prime pattern is borrowed from the open-source `gridstatus.io` library, which uses the same approach for several ISO-NE endpoints. The cookie is short-lived (ses-

sion duration only); we re-prime on each Python process start. There is no rate-limit on the documented endpoint as of writing; we have observed no 429s in production.

C Datasheet for the public dataset

Following Gebru et al. (8).

Motivation. The dataset supports day-ahead 24-hour electricity demand forecasting for the eight ISO New England load zones. It pairs hourly per-zone demand (from ISO-NE) with hourly weather rasters (from NOAA HRRR) over a common UTC time index, plus deterministic 44-d calendar one-hots. The dataset is intended for ML/forecasting research; it is not certified for production grid operations.

Composition. Each sample is a 48-hour window with $S = 24$ history hours and $H = 24$ future hours. Per timestep: a $(450, 449, 7)$ HRRR weather raster, an 8-d per-zone demand vector, and a 44-d calendar one-hot. Per sample: history+future weather (48 timesteps), history demand (24 timesteps), future demand target (24 timesteps), history+future calendar (48 timesteps). Calendar years 2019–2023 yield $\sim$26 K samples after dropping any timesteps with missing weather or demand.

Collection process. ISO-NE per-zone demand is published through the public Energy/Load-/Demand reports portal (11) as monthly CSVs by zone, plus a 5-minute live feed (Appendix B). HRRR weather is available from NOAA's S3 mirror `s3://noaa-hrrr-bdp-pds/` (2, 18); we use the f00 analysis product. Calendar features are computed deterministically from UTC timestamps. No PII; no human subjects; no sensitive demographic data.

Preprocessing. (a) Demand: monthly per-zone CSVs are joined into a 2019–2023 hourly long table; UTC alignment uses pandas with US/Eastern timezone localization. (b) Weather: HRRR f00 analyses are fetched via Herbie at hourly cadence, restricted to the New England bbox, and bilinearly regridded to a regular 450×449 lat/lon grid. (c) Calendar: 44-d one-hot via custom encoder. (d) Splits: train (2019–2020 baseline / 2019–2021 encoder-decoder) / val (2021 baseline / 2022 encoder-decoder) / test (last 2 days of 2022). All MAPE numbers in this paper are computed on the test slice or on the live deployment-period rolling window.

Uses. Day-ahead 24-hour per-zone demand forecasting; comparison of multimodal vs foundation-model time-series forecasters; distribution-shift case studies for deployed forecasting systems. *Not* certified for grid operations or market dispatch.

Distribution. The full pre-processed dataset is reproducible from public sources via `scripts/data_preparation` `fetch_iso_ne_demand.py` (per-zone monthly CSVs; downloads from ISO-NE) + `fetch_hrrr_weather.py` (per-hour HRRR analyses; downloads from NOAA AWS S3 via Herbie). The deployed live demo additionally consumes the 5-minute zonal live feed (Appendix B).

Maintenance. The auxiliary data repository (`new-england-real-time-power-predict-data`) refreshes the rolling 7-day backtest daily at 04:00 UTC via GitHub Actions. The training-time CSVs

and HRRR cache are pinned to the snapshot used for this paper. Source code is at `real-time-power-predict`.

PII. None. The data is aggregate per-zone load and gridded weather analyses; no individual customer data is exposed.

D Reproducibility

Table 7: Hyperparameters for the baseline (§4) and the encoder-decoder variants (§5).

	Baseline	Encoder-decoder (both variants)
History length S	24	24
Future length H	24	24
Spatial token grid	8×8 (P=64)	8×8 (P=64)
Embedding dim D	128	128
Encoder layers	4	4
Decoder layers	—	2
Attention heads	4	4
Dropout	0.1	0.1
Optimizer	AdamW (15)	AdamW
Initial LR	3×10^{-4}	3×10^{-4}
Weight decay	10^{-2}	10^{-2}
Schedule	cosine warm-restarts (14)	cosine warm-restarts
Batch size	16	16
Epochs (planned / actual)	14 / 14	24 / 12 (v1), 24 / 13 (v2)
Loss	MSE in z-score space	MSE in z-score space
Mixed precision	PyTorch AMP fp16	PyTorch AMP fp16
Random seeds	not set	not set

Hyperparameters.

Table 8: Validation MAPE per training epoch (lower is better). Encoder-decoder variants used a chained `-resume` workflow that reset the LR scheduler at every chain boundary; the kink at epoch 7 is visible in Figure 2.

Epoch	Baseline (val 2021)	v1 (val 2022)	v2 (val 2022, undertrained)
0	11.22 %	10.08 %	—
4	9.16 %	8.70 %	—
6	8.76 %	**8.63 %** ($\star$)	—
8	—	9.95 %	8.95 % (LR reset to 1e-3)
13	**6.92 %**	10.50 %	**8.72 %** ($\star$)

($\star$) = best-of-run; both encoder-decoder variants stalled after the LR-scheduler reset.

Per-epoch validation MAPE trajectory.

Random seeds caveat. The training pipeline does not set `torch.manual_seed` / `np.random.seed` / `torch.use_deterministic_algorithms`. The reported test MAPE numbers (5.24 % baseline, 6.82 % encoder-decoder, 4.33 % / 4.21 % ensembles) are pinned to the specific checkpoints in `runs/cnn_transformer_baseline/checkpoints/best.pt` and `runs/cnn_encoder_decoder/checkpoints/best` both bundled with the public release; they are not bit-reproducible across re-training runs.

Validation script. The deployment-validation experiment in §9 is executable as a single Python script:

```
$ python scripts/validation/reproduce_dec30_2022.py
# fetches HRRR for 2022-12-29 (history) + 2022-12-30 (future analyses)
# runs deployed pipeline, compares to the cluster's stored prediction
# expected output:
#   live MAPE       : 6.41 %
#   cluster MAPE    : 6.54 %
#   max abs diff    : 94 MW (3.6 %)
```

Compute. Training was done on NVIDIA A100-80 GB / A100-40 GB / P100 GPUs. Each baseline epoch takes ∼45 minutes; the full 14-epoch training is ∼10 hours of GPU time. Encoder-decoder epochs are slightly slower (∼60 min/epoch) due to the additional decoder stack. Inference is CPU-feasible: HuggingFace Spaces' free CPU tier (2 vCPU, 16 GB RAM) serves a forward pass in ∼5 s when warm.

Build instructions. The arXiv source is at `report/arxiv/` in the public code repository. To rebuild:

```
$ cd report/arxiv
$ latexmk -xelatex -interaction=nonstopmode -halt-on-error paper.tex
```

The figures shipped with the paper are deterministically regenerable via the scripts in `scripts/figures/`; the live-deployment numbers in Table 6 and Figures 10–11 pull from the data-repo URL and will reflect the most recent cron build at compile time. The version of the paper distributed with this repository pins to the 2026-05-05 cron snapshot; subsequent builds will pick up newer windows.

E Attention extraction procedure

The four orientation-correctness checks used by `scripts/attention_maps.py` (§7):

1. **Slice direction**. Rows of the sliced attention matrix are future-tabular query positions $\mathcal{Q}_{\text{fut}}$; columns are history-spatial key positions $\mathcal{K}_{\text{hist}}$ (Equation 9).
2. **Row-major reshape**. The 64-cell column dimension is reshaped to $(8, 8)$ row-major, matching `WeatherCNN.forward`'s `x.flatten(2)` layout: cell index $i \to$ (row$=i \div 8$, col$=i \bmod 8$).
3. **Compass orientation**. `matplotlib.imshow(..., origin='upper')` so row 0 (north) is at the top, col 0 (west) is on the left.
4. **Aggregate east-vs-west sanity**. Equation 10; if violated, the script emits a `[WARN]` and the figure is flagged in the report rather than aborted.

The slice operation.

$$A^{(\ell,h)}_{\text{slice}} \;=\; A^{(\ell,h)}[\mathcal{Q}_{\text{fut}}, \mathcal{K}_{\text{hist}}] \;\in\; \mathbb{R}^{|\mathcal{Q}_{\text{fut}}|\times|\mathcal{K}_{\text{hist}}|} \;=\; \mathbb{R}^{24\times 1536}, \tag{9}$$

with $|\mathcal{Q}_{\text{fut}}| = 24$ future-tabular query positions and $|\mathcal{K}_{\text{hist}}| = S \times P = 24 \times 64 = 1536$ history-spatial key positions. We then sum over the history time axis to get a 24×64 map of future-time vs spatial position, and reshape the spatial axis to 8×8.

The east-vs-west sanity check.

$$\sum_{r=0}^{7}\sum_{c=4}^{7} A_{\text{agg}}[r,c] \;>\; \sum_{r=0}^{7}\sum_{c=0}^{3} A_{\text{agg}}[r,c], \tag{10}$$

i.e., aggregate attention mass on the eastern (cols 4–7) half of the bbox should exceed the western half because dense load (Boston metro, southern Massachusetts) sits in the eastern half. Observed values: east ≈ 0.0070, west ≈ 0.0050 — check passes.

Diagnostic days. The five hand-picked days used in §7 are 2022-04-15 (mild), 2022-07-21 (heat-wave), 2022-08-04 (extreme), 2022-12-23 (cold-snap), and 2022-12-25 (holiday). Drawn from the broader 2022 calendar; the held-out test slice 2022-12-30/31 is too narrow to span all five regimes.

F AI-tool disclosure

The author used *Claude Code* (Anthropic; models: Opus 4.7 and Sonnet 4.6) as a coding and writing assistant throughout this project. AI was used for: (i) code scaffolding (training-loop boilerplate, SLURM scripts, Gradio app stub, the rolling-backtest cron, the reproducibility validation script); (ii) figure-rendering code (`matplotlib` generators in `scripts/figures/` and the TikZ source for the architecture diagrams); (iii) LaTeX boilerplate (`preamble.tex`, citation expansion in `refs.bib`); (iv) prose copy-editing on this paper.

AI was *not* used for: research design, model architecture choices, training-result interpretation, the LR-scheduler-reset-bug diagnosis, the four geographic-orientation correctness checks in §7, or any final scientific claim. All design decisions, all numerical results, and all interpretation in this paper reflect the author's own judgment and responsibility.